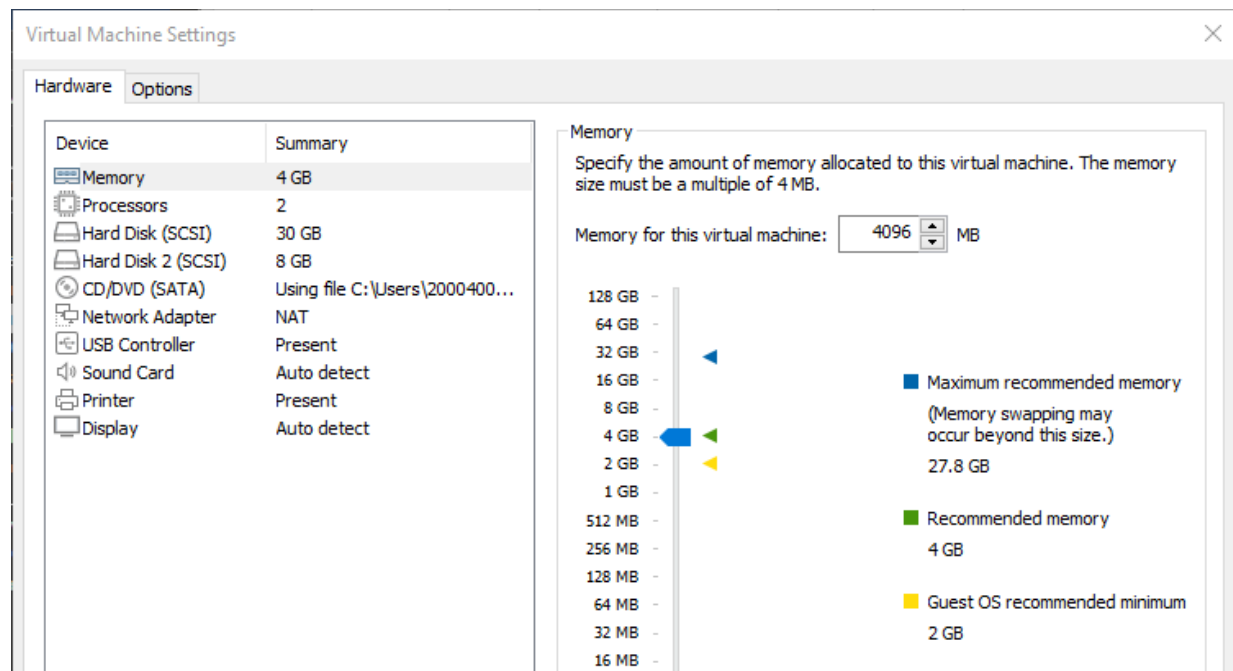
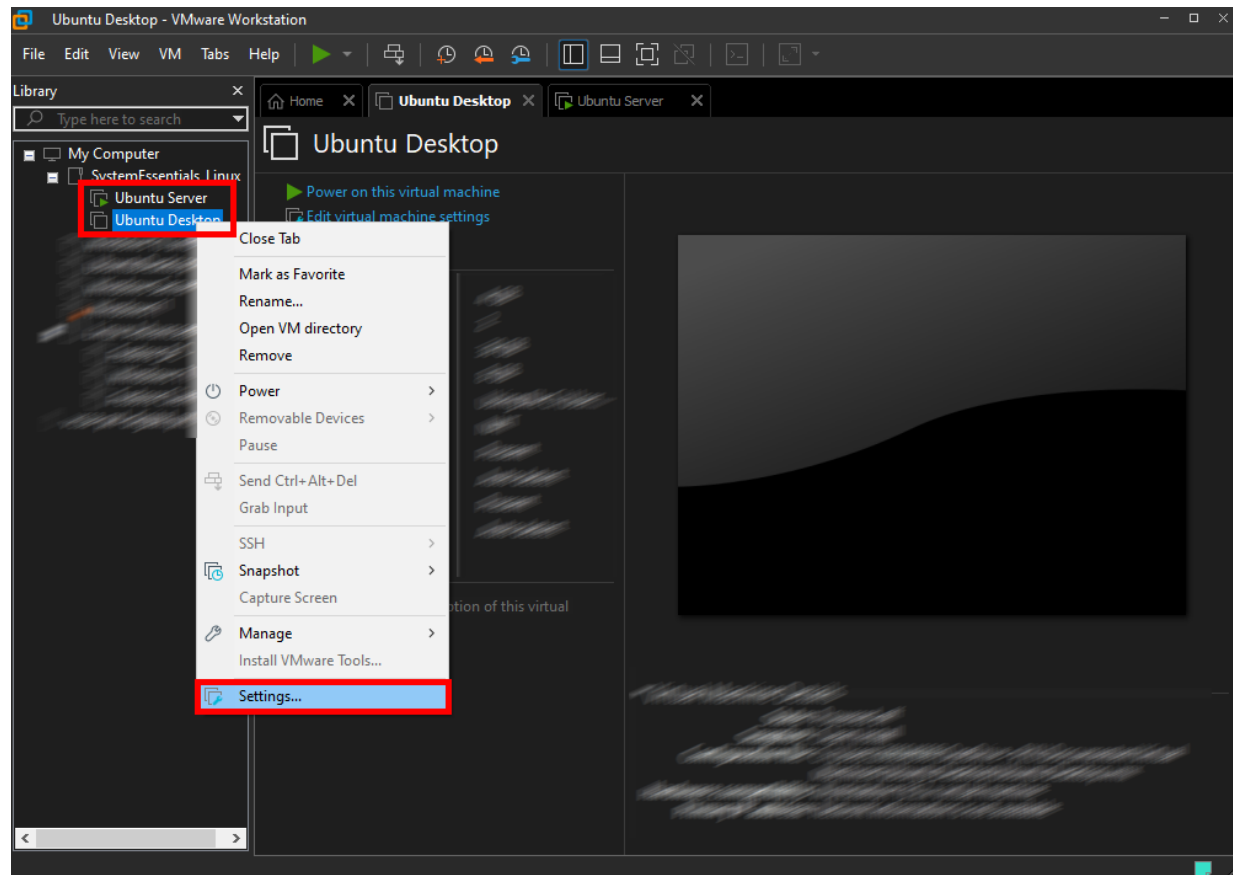


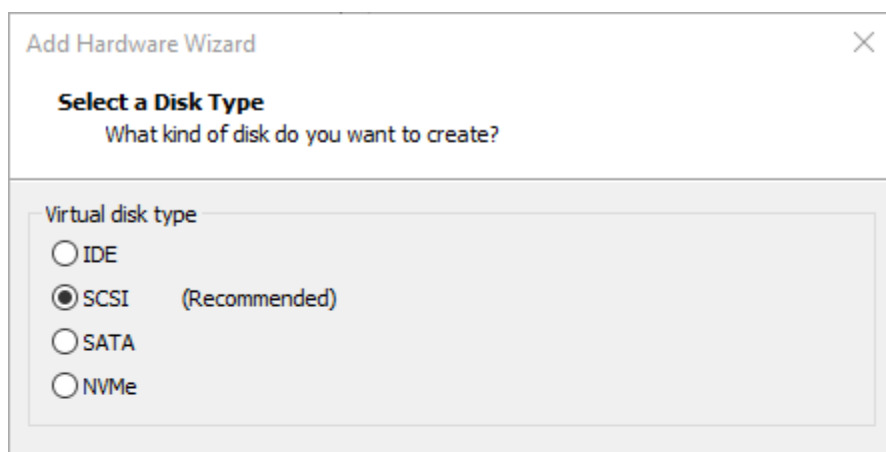
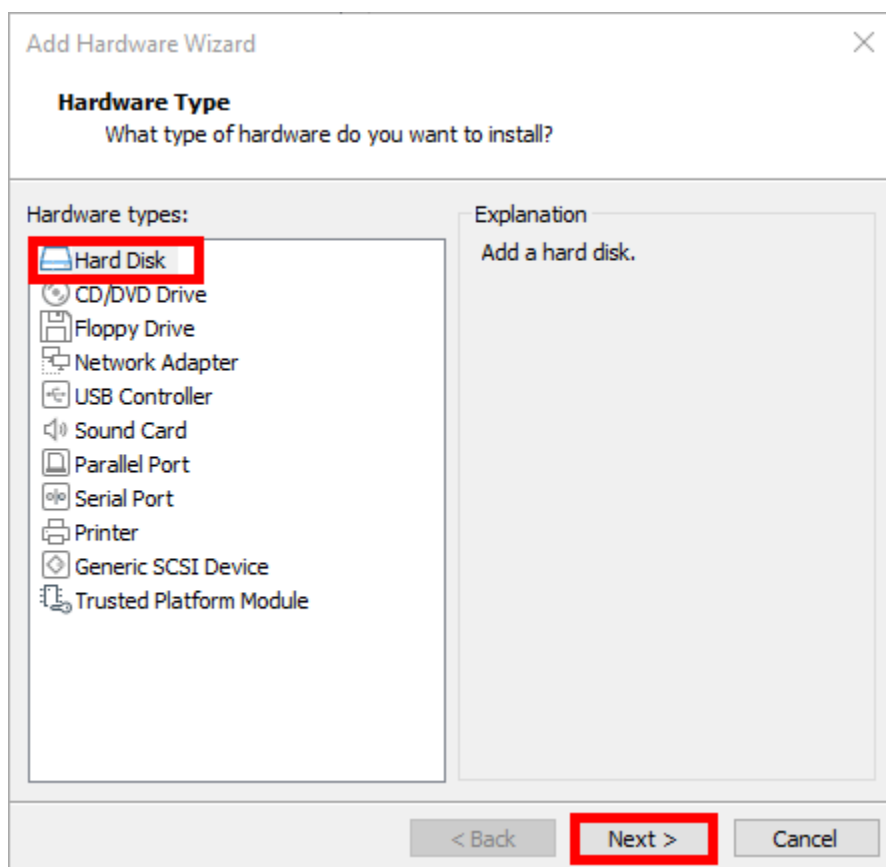
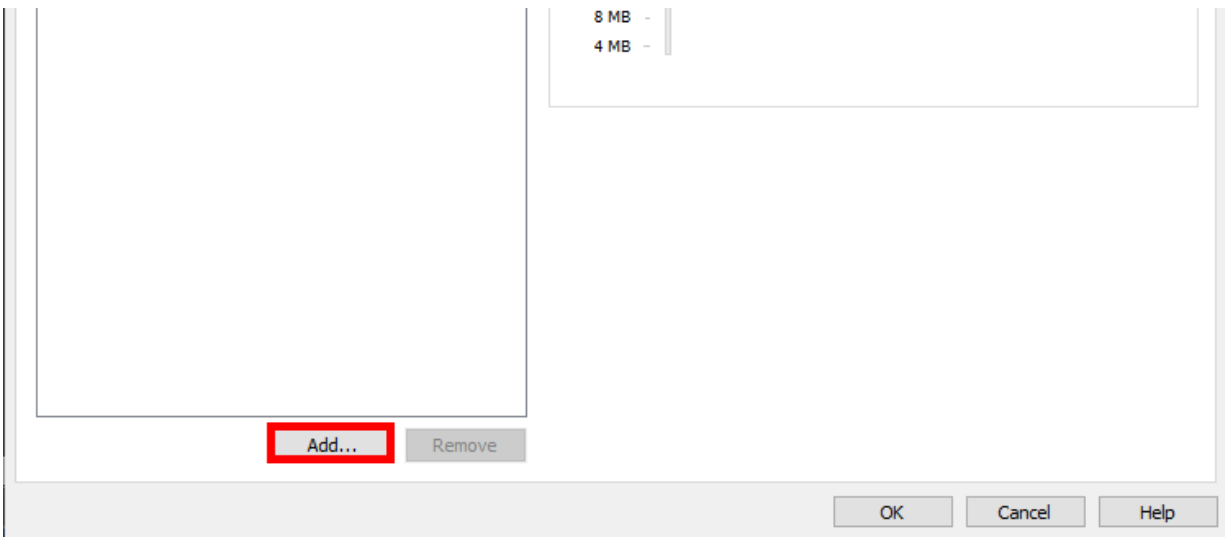
Managing drives and filesystems

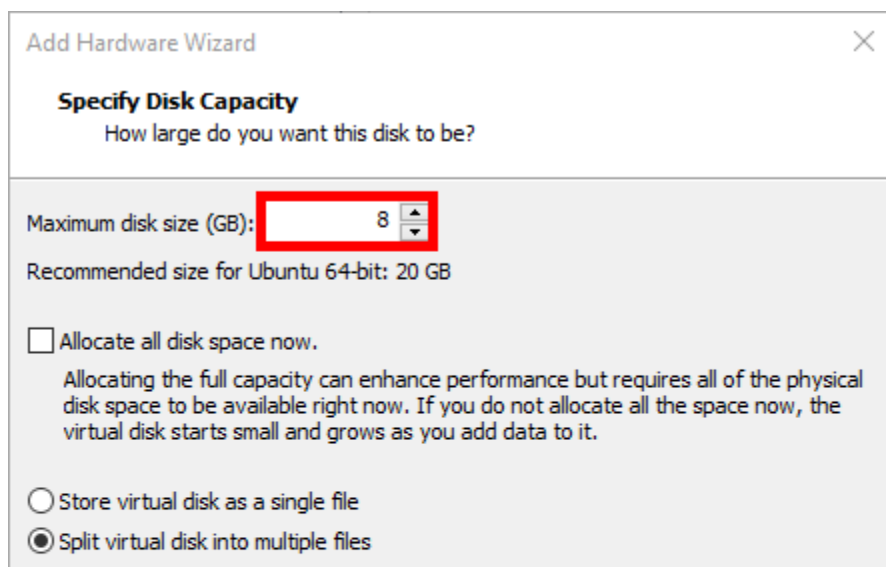
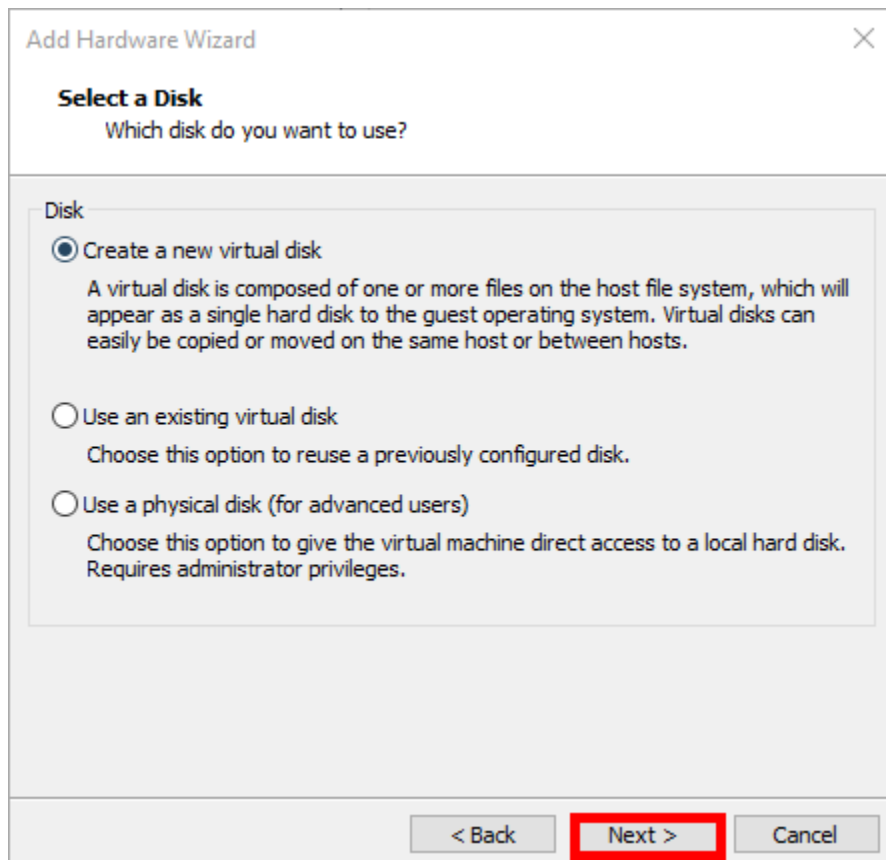

Every hard drive that is used in our system is not just plug and play, to use it we need to partition it first. Partitioning means cutting in parts, these parts then receive a certain structure your computer can work with also called a file system. The basics of data storage is almost the same for most operating systems nowadays. When installing an operating system the drive we use, HDD, SSD or NVME, gets partitioned and each partition gets formatted with a certain file system. In Linux some partitions are formatted in a special way, namely the swap partition and Linux Logical Volume Manager (LVM). In our system, HDD's, SSD's and NVME's are used for permanent storage. RAM and swap are used for temporary storage. For example when running a command, the command gets copied from your drive into RAM to execute it more quickly. This is because our RAM can be read faster than drives by our CPU. However RAM has a smaller capacity due to it's cost and RAM gets deleted every time our computer is shut down. When our RAM runs out of memory, we can use the swap partition to temporarily hold the RAM parts that are not in use and load other data into RAM. The other special partition is the Linux LVM. This logical volume lets us create pools of storage called volume groups. These groups give more flexibility to enlarge or shrink logical volumes than normal drive partitions offer. In our Linux distribution at least 1 partition is necessary, the root partition or "/". But most of the time multiple partitions are created, for example the directories /home, /var or /tmp. Each of these directories get mounted to a mount point under the / partition. When adding files or folders to this mount point, the files and folders get saved on this different partition. These different partitions and mounting of them happens automatically so it's not visible for the end user. Each of our disk partitions gets a device name when booting Linux, for example sda2. An entry in the file /etc/fstab tells Linux where to mount each partition, which occurs at startup. Note the difference with Windows that everything is mounted under the root (/) and not under a drive letter like C:, D:, ... Some drives are automatically mounted to our file system when inserting a removable media. For example, a CD-ROM can get mounted under /media/cdrom or `/run/media/<username>/<cdrom name>` . When this doesn't occur automatically, an administrator should

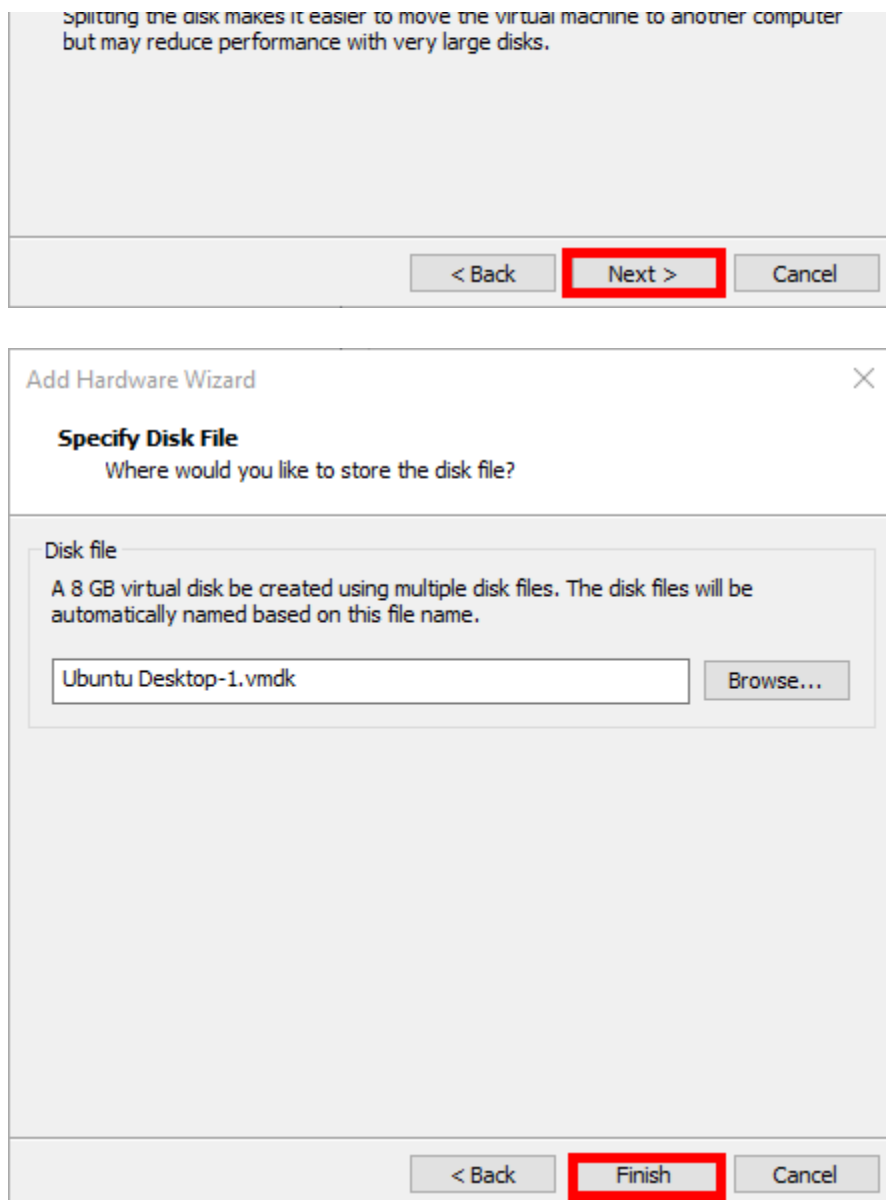
create the mount point to a folder of his/her choice. Linux is able to work with VFAT, used in USB sticks, handy when exchanging files with a windows system. It also has kernel support for New Technology File System (NTFS), but there are often additional drivers needed to load NTFS.

To start this chapter, we'll need to add drives to our virtual machine first.









Reboot your Virtual Machine if it was still running, otherwise start it and the drive will be recognised by the kernel and ready to be initialised. Be sure to take a snapshot before continuing this chapter and to double check the steps taken. Be careful and do not repartition the drive on which Linux is installed! A bad entry in the `/etc/fstab` file could also result in an unbootable Linux.

Understanding partition tables and disk partitions

Traditionally MBR partition tables were used to save the size and layout of partitions. In Linux a lot of tools are available to do this. But nowadays the new standard Global Unique Identifiers, GUID, partition tables are used. Your

computer needs to use UEFI for this standard. This change occurred because of the limits of MBR. MBR partitions could be a maximum of 2TB with a maximum of 4 primary partitions or 3 primary and 1 extended partition (which can hold multiple logical partitions). The GUID partitions can get to a maximum of 9,4 ZB (Zettabytes or 10^{21} bytes) with a maximum of 128 partitions. We can use the commands `fdisk` or `gdisk` to partition a drive. `gdisk` gives the possibility to create larger partitions than `fdisk`, other sub commands to create, delete or change partitions are more or less the same.

bash

```
student@linux-ess:~$ sudo fdisk -l /dev/sdb
Disk /dev/sdb: 8 GiB, 8589934592 bytes, 16777216 sectors
Disk model: VMware Virtual S
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
student@linux-ess:~$
student@linux-ess:~$ sudo gdisk -l /dev/sdb
GPT fdisk (gdisk) version 1.0.8

Partition table scan:
  MBR: protective
  BSD: not present
  APM: not present
  GPT: present

Found valid GPT with protective MBR; using GPT.
Disk /dev/sdb: 16777216 sectors, 8.0 GiB
Model: VMware Virtual S
Sector size (logical/physical): 512/512 bytes
Disk identifier (GUID): 4B226FAA-BFB1-4A6F-9F76-CA70124E6336
Partition table holds up to 128 entries
Main partition table begins at sector 2 and ends at sector 33
First usable sector is 34, last usable sector is 16777182
Partitions will be aligned on 2048-sector boundaries
Total free space is 16777149 sectors (8.0 GiB)
```

Number	Start (sector)	End (sector)	Size	Code	Name
--------	----------------	--------------	------	------	------

Every SCSI, SATA or USB device gets represented by sd? (sda, sdb, sdc, ...). With MBR these devices can have a maximum of 16 subdivisions (sdc, sdc1 -> sdc15) this means there is a maximum of 15 partitions with MBR. When using GPT there is a maximum of 128 partitions (sdd, sdd1 -> sdd127). With MBR, a drive can have 4 (primary) partitions maximum. If you need more than 4 partitions, you'll need to use extended partition(s) with logical partitions. The first drive mostly shows as /dev/sda. As said before at least one partition is created when installing Linux, this partition is used as a Linux LVM physical partition where other logical partitions can be created.

bash

```
student@linux-ess:~$ sudo fdisk -l /dev/sda
[sudo] password for student:
Disk /dev/sda: 20 GiB, 21474836480 bytes, 41943040 sectors
Disk model: VMware Virtual S
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disklabel type: gpt
Disk identifier: A9E92175-1433-43DC-968D-95C5E18A2105

Device        Start      End  Sectors  Size Type
/dev/sda1      2048      4095    2048    1M BIOS boot
/dev/sda2      4096  3719167  3715072  1.8G Linux filesystem
/dev/sda3  3719168 41940991 38221824 18.2G Linux filesystem
student@linux-ess:~$
student@linux-ess:~$ sudo gdisk -l /dev/sda
GPT fdisk (gdisk) version 1.0.8
```

Partition table scan:

```
MBR: protective
BSD: not present
APM: not present
GPT: present
```

Found valid GPT with protective MBR; using GPT.

```
Disk /dev/sda: 41943040 sectors, 20.0 GiB
Model: VMware Virtual S
Sector size (logical/physical): 512/512 bytes
```

Disk identifier (GUID): A9E92175-1433-43DC-968D-95C5E18A2105

Partition table holds up to 128 entries

Main partition table begins at sector 2 and ends at sector 33

First usable sector is 34, last usable sector is 41943006

Partitions will be aligned on 2048-sector boundaries

Total free space is 4029 sectors (2.0 MiB)

Number	Start (sector)	End (sector)	Size	Code	Name
1	2048	4095	1024.0 KiB	EF02	
2	4096	3719167	1.8 GiB	8300	
3	3719168	41940991	18.2 GiB	8300	

Looking at our sda drive we see it partitioned into /boot of +- 1GB. The * indicates this partition is bootable. The rest of the drive is a physical LVM partition. This one is used to create logical volumes. With lsblk we see all available drives and their partitions.

bash

```
student@linux-ess:~$ lsblk
```

\$NAME	MAJ:MIN	RM	SIZE	RO	TYPE	MOUNTPPOINTS
loop0	7:0	0	63.2M	1	loop	/snap/core20/1623
loop1	7:1	0	48M	1	loop	/snap/snapd/16778
loop2	7:2	0	79.9M	1	loop	/snap/lxd/22923
loop3	7:3	0	62M	1	loop	/snap/core20/1587
loop4	7:4	0	103M	1	loop	/snap/lxd/23541
loop5	7:5	0	47M	1	loop	/snap/snapd/16292
sda	8:0	0	20G	0	disk	
└─sda1	8:1	0	1M	0	part	
└─sda2	8:2	0	1.8G	0	part	/boot
└─sda3	8:3	0	18.2G	0	part	
└─ubuntu--vg-ubuntu--lv	253:0	0	10G	0	lvm	/
sdb	8:16	0	8G	0	disk	
sr0	11:0	1	1.4G	0	rom	

Adding a disk with one partition

Next up, we'll partition the new drive and install a file system. Afterwards we'll be able to mount the new partition to a folder in Linux. The easiest method is to use the full drive space for one partition. It is possible to add multiple partitions, you'll need to add a file system to each partition afterwards and each partition needs to be mounted separately. If a partition of a drive is mounted, like a USB stick, and you want to repartition it, you'll need to unmount it first. We'll now start with creating a partition and installing a file system on our extra drive. If a mistake is made while working with the `fdisk/gdisk` command, finish the current operation and press `q` to quit without saving.

First up, check the device name of our newly added drive by using the `lsblk` command

```
bash
student@linux-ess:~$ lsblk
```

\$NAME	MAJ:MIN	RM	SIZE	RO	TYPE	MOUNTPOINTS
loop0	7:0	0	63.2M	1	loop	/snap/core20/1623
loop1	7:1	0	48M	1	loop	/snap/snapd/16778
loop2	7:2	0	79.9M	1	loop	/snap/lxd/22923
loop3	7:3	0	62M	1	loop	/snap/core20/1587
loop4	7:4	0	103M	1	loop	/snap/lxd/23541
loop5	7:5	0	47M	1	loop	/snap/snapd/16292
sda	8:0	0	20G	0	disk	
└─sda1	8:1	0	1M	0	part	
└─sda2	8:2	0	1.8G	0	part	/boot
└─sda3	8:3	0	18.2G	0	part	
└─ubuntu--vg-ubuntu--lv	253:0	0	10G	0	lvm	/
sdb	8:16	0	8G	0	disk	
sr0	11:0	1	1.4G	0	rom	

```
# 0
```

In this case we'll need to use `/dev/sdb`. Now start by using the `fdisk` or `gdisk` command and the drive we want to use.

```
bash
student@linux-ess:~$ sudo fdisk /dev/sdb
```

```
Welcome to fdisk (util-linux 2.37.2).
Changes will remain in memory only, until you decide to write them.
Be careful before using the write command.
```

```
Device does not contain a recognized partition table.  
Created a new DOS disklabel with disk identifier 0x8b8ca071.
```

```
Command (m for help):
```

```
bash
```

```
student@linux-ess:~$ sudo gdisk /dev/sdb  
GPT fdisk (gdisk) version 1.0.8
```

```
Partition table scan:
```

```
  MBR: protective  
  BSD: not present  
  APM: not present  
  GPT: present
```

```
Found valid GPT with protective MBR; using GPT.
```

```
Command (? for help):
```

We'll first check if our drive isn't already formatted. Press p to check for partitions.

```
bash
```

```
# FDISK  
Command (m for help): p  
Disk /dev/sdb: 8 GiB, 8589934592 bytes, 16777216 sectors  
Disk model: VMware Virtual S  
Units: sectors of 1 * 512 = 512 bytes  
Sector size (logical/physical): 512 bytes / 512 bytes  
I/O size (minimum/optimal): 512 bytes / 512 bytes  
Disklabel type: dos  
Disk identifier: 0x8b8ca071
```

```
Command (m for help):
```

bash

```
# GDISK
Command (? for help): p
Disk /dev/sdb: 16777216 sectors, 8.0 GiB
Model: VMware Virtual S
Sector size (logical/physical): 512/512 bytes
Disk identifier (GUID): 4B226FAA-BFB1-4A6F-9F76-CA70124E6336
Partition table holds up to 128 entries
Main partition table begins at sector 2 and ends at sector 33
First usable sector is 34, last usable sector is 16777182
Partitions will be aligned on 2048-sector boundaries
Total free space is 16777149 sectors (8.0 GiB)

Number  Start (sector)    End (sector)  Size      Code  Name

Command (? for help):
```

If a partition would be present, we'll need to delete it before going on. You could do this by pressing d. It will then tell you what partition is selected, pressing enter will delete this partition. We are now ready to create a new partition, do this by pressing n. We'll now be prompted to choose a primary (p), or extended (e) partition with fdisk, gdisk doesn't ask for this. As this is our first partition, choose the primary partition by pressing p. Afterwards enter the partition number, again as this is our first partition, we'll use number 1. Do this by pressing 1 and pressing enter. Next we'll be prompted where the first sector should start, we'll use the default value, so just press enter. Now it's time to set the size of the partition, you could enter a number, but we'll use the full size of our drive. Enter the number of the last sector or just press enter. Now we can check our newly created partition by pressing p. If everything is as expected, press w to write or save the changes to the partition table.

bash

```
# FDISK
Command (m for help): n
Partition type
   p   primary (0 primary, 0 extended, 4 free)
   e   extended (container for logical partitions)
```

```
Select (default p): p
Partition number (1-4, default 1): 1
First sector (2048-16777215, default 2048):
Last sector, +/-sectors or +/-size{K,M,G,T,P} (2048-16777215, default 167
```

Created a new partition 1 of type 'Linux' and of size 8 GiB.

```
Command (m for help): p
Disk /dev/sdb: 8 GiB, 8589934592 bytes, 16777216 sectors
Disk model: VMware Virtual S
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disklabel type: dos
Disk identifier: 0x8b8ca071
```

Device	Boot	Start	End	Sectors	Size	Id	Type
/dev/sdb1		2048	16777215	16775168	8G	83	Linux

```
Command (m for help): w
The partition table has been altered.
Calling ioctl() to re-read partition table.
Syncing disks.
```

bash

```
# GDISK
Command (? for help): n
Partition number (1-128, default 1): 1
First sector (34-16777182, default = 2048) or {+-}size{KMGTP}:
Last sector (2048-16777182, default = 16777182) or {+-}size{KMGTP}:
Current type is 8300 (Linux filesystem)
Hex code or GUID (L to show codes, Enter = 8300):
Changed type of partition to 'Linux filesystem'
```

```
Command (? for help): p
Disk /dev/sdb: 16777216 sectors, 8.0 GiB
Model: VMware Virtual S
Sector size (logical/physical): 512/512 bytes
```

```
Disk identifier (GUID): 4B226FAA-BFB1-4A6F-9F76-CA70124E6336
Partition table holds up to 128 entries
Main partition table begins at sector 2 and ends at sector 33
First usable sector is 34, last usable sector is 16777182
Partitions will be aligned on 2048-sector boundaries
Total free space is 2014 sectors (1007.0 KiB)
```

Number	Start (sector)	End (sector)	Size	Code	Name
1	2048	16777182	8.0 GiB	8300	Linux filesystem

```
Command (? for help): w
```

```
Final checks complete. About to write GPT data. THIS WILL OVERWRITE EXIST
PARTITIONS!!
```

```
Do you want to proceed? (Y/N): Y
```

```
OK; writing new GUID partition table (GPT) to /dev/sdb.
```

```
Warning: The kernel is still using the old partition table.
```

```
The new table will be used at the next reboot or after you
run partprobe(8) or kpartx(8)
```

```
The operation has completed successfully.
```

If this should fail its probably because the drive was still mounted. Unmount the drive and use the partprobe command to save the changes to the partition table. If that still fails, reboot your computer and try again.

bash

```
student@linux-ess:~$ sudo partprobe /dev/sdb
```

The new partition still is not ready for use. We need to install a file system on it as mentioned before. Use the mkfs command to accomplish this. The standard mkfs command creates a ext2 file system, normally you would like a journaling system like ext3/ ext4. Change the mkfs command with parameter -t to get this. Another option is to use the mkfs.ext4 command, this is a shorter option. We could also use a xfs-file system.

bash

```
student@linux-ess:~$ sudo mkfs -t ext4 /dev/sdb1
mke2fs 1.46.5 (30-Dec-2021)
Creating filesystem with 2096896 4k blocks and 524288 inodes
Filesystem UUID: f2fee344-9305-46d7-9942-339fd434dd8d
Superblock backups stored on blocks:
    32768, 98304, 163840, 229376, 294912, 819200, 884736, 1605632
```

```
Allocating group tables: done
Writing inode tables: done
Creating journal (16384 blocks): done
Writing superblocks and filesystem accounting information: done
```

```
student@linux-ess:~$ sudo mkfs.ext4 /dev/sdb1
mke2fs 1.46.5 (30-Dec-2021)
Creating filesystem with 2096896 4k blocks and 524288 inodes
Filesystem UUID: f2fee344-9305-46d7-9942-339fd434dd8d
Superblock backups stored on blocks:
    32768, 98304, 163840, 229376, 294912, 819200, 884736, 1605632
```

```
Allocating group tables: done
Writing inode tables: done
Creating journal (16384 blocks): done
Writing superblocks and filesystem accounting information: done
```

```
student@linux-ess:~$ sudo mkfs -t xfs /dev/sdb1
meta-data=/dev/sdb1          isize=512    agcount=4, agsize=524224 bl
        =                   sectsz=512   attr=2,    projid32bit=1
        =                   crc=1        finobt=1, sparse=1, rmapbt=
        =                   reflink=1    bigtime=0 inobtcount=0
data      =                   bsize=4096   blocks=2096896, imaxpct=25
        =                   sunit=0      swidth=0 blks
naming    =version 2          bsize=4096   ascii-ci=0, ftype=1
log       =internal log      bsize=4096   blocks=2560, version=2
        =                   sectsz=512   sunit=0 blks, lazy-count=1
realtime  =none              extsz=4096   blocks=0,   rtextents=0
```

```
student@linux-ess:~$ sudo mkfs.xfs /dev/sdb1
meta-data=/dev/sdb1          isize=512    agcount=4, agsize=524224 bl
```

```

= sectsz=512 attr=2, projid32bit=1
= crc=1 finobt=1, sparse=1, rmapbt=
= reflink=1 bigtime=0 inobtcount=0
data = bsize=4096 blocks=2096896, imaxpct=25
= sunit=0 swidth=0 blks
naming =version 2 bsize=4096 ascii-ci=0, ftype=1
log =internal log bsize=4096 blocks=2560, version=2
= sectsz=512 sunit=0 blks, lazy-count=1
realtime =none extsz=4096 blocks=0, rtextents=0

```

Now that we created a partition and created a file system on it, it's time to mount our newly created partition to a mount point. Use the mount command to do this.

bash

```

student@linux-ess:~$ sudo mkdir /mnt/test
student@linux-ess:~$ sudo mount /dev/sdb1 /mnt/test
student@linux-ess:~$ sudo df -h /mnt/test
Filesystem      Size  Used Avail Use% Mounted on
/dev/sdb1        7.8G   24K   7.4G   1% /mnt/test
student@linux-ess:~$ sudo mount | grep sdb1
/dev/sdb1 on /mnt/test type ext4 (rw,relatime)

```

With the df command we see that /dev/sdb1 is mounted to the folder /mnt/test. We can also see it provides 7.8GB of memory. The mount command shows all mounted partitions, with grep we filter to see only our new partition. When the partition is no longer in needed, we can unmount it with the umount command

bash

```

student@linux-ess:~$ sudo umount /dev/sdb1

```

With the mount command, the partition is temporarily mounted and gets unmounted when the pc gets rebooted. When we want the partition to be remounted on startup we'll need to add an entry for it in the /etc/fstab file.

bash

```
student@linux-ess:~$ sudo nano /etc/fstab
# /etc/fstab: static file system information.
#
# Use 'blkid' to print the universally unique identifier for a
# device; this may be used with UUID= as a more robust way to name device
# that works even if disks are added and removed. See fstab(5).
#
# <file system> <mount point>   <type>  <options>          <dump>  <pass>
# / was on /dev/ubuntu-vg/ubuntu-lv during curtin installation
/dev/disk/by-id/dm-uuid-LVM-cPT0cPGozhK91HrSeQ2ByYSCtWdNqE7Co7jRC002XWz9d
# /boot was on /dev/sda2 during curtin installation
/dev/disk/by-uuid/0b189ed9-2d70-4955-9e36-66ad45dbbee7 /boot ext4 default
/dev/sdb1 /mnt/test ext4 defaults 0 2
```

In this example `/dev/sdb1` gets mounted to the folder `/mnt/test` with `ext4` as its file system. The word `defaults` means that it gets mounted with default options (`rw`, `auto`, ...). The two numbers at the end stand for: the first one tells the system it does not need to make backup files of this file system with the help of the `dump` command. This command is barely used nowadays, but the number field is still available. The second one indicates the order on which the system checks this mount. This value is 1 for the root file system and 2 for others. If you do not want to check this partition at boot time set this field to 0. Now that this partition is added to the `/etc/fstab` file, it will be mounted every time your system boots.

adding a disk with multiple partitions

We'll now start again with our drive and try to create multiple partitions. The partitions we want are:

- `sdb1` and `sdb2`: 500 MB
- `sdb3`: 300MB
- `sdb5`: 350 MB
- `sdb6`: 400 MB

Sdb4, which is missing in our list, is an extended partition when using fdisk, which takes all remaining disk space (after sdb1, sdb2 and sdb3).

sdb5 and sdb6 are logical partitions within the extended partition (sdb4).

When using gdisk this is not needed, but we will use the same numbers to keep consistency.

First, unmount all mounted partitions of this drive. And if they were mounted automatically with an entry in the /etc/fstab file, you'll also need to remove those entries from /etc/fstab.

bash

```
student@linux-ess:~$ sudo umount /dev/sdb1
student@linux-ess:~$ sudo nano /etc/fstab      #to remove the sdb1 entry
```

Now we go back to our fdisk or gdisk command.

bash

```
student@linux-ess:~$ sudo fdisk /dev/sdb
```

Welcome to **fdisk** (util-linux 2.37.2).

Changes will remain **in** memory only, **until** you decide to **write** them.

Be careful before using the **write** command.

Command (m **for** help):

bash

```
student@linux-ess:~$ sudo gdisk /dev/sdb
```

GPT **fdisk** (gdisk) version 1.0.8

Partition table scan:

MBR: protective

BSD: not present

APM: not present

GPT: present

Found valid GPT with protective MBR; using GPT.

Command (? for help):

Check for existing partitions, we know there is one, by pressing p. We'll delete this partition by entering d, checking the selected partition and if correct, pressing enter.

bash

```
# FDISK
Command (m for help): p
Disk /dev/sdb: 8 GiB, 8589934592 bytes, 16777216 sectors
Disk model: VMware Virtual S
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disklabel type: dos
Disk identifier: 0x8b8ca071
```

Device	Boot	Start	End	Sectors	Size	Id	Type
/dev/sdb1		2048	16777215	16775168	8G	83	Linux

```
Command (m for help): d
Selected partition 1
Partition 1 has been deleted.
```

Command (m for help):

bash

```
# GDISK
Command (? for help): p
Disk /dev/sdb: 16777216 sectors, 8.0 GiB
Model: VMware Virtual S
Sector size (logical/physical): 512/512 bytes
Disk identifier (GUID): 4B226FAA-BFB1-4A6F-9F76-CA70124E6336
Partition table holds up to 128 entries
Main partition table begins at sector 2 and ends at sector 33
First usable sector is 34, last usable sector is 16777182
Partitions will be aligned on 2048-sector boundaries
```

Total free space is 2014 sectors (1007.0 KiB)

Number	Start (sector)	End (sector)	Size	Code	Name
1	2048	16777182	8.0 GiB	8300	Linux filesystem

Command (? for help): d

Using 1

Command (? for help):

Now, we'll create new partitions by entering n, for sdb1, sdb2 and sdb3 we'll use primary partitions. So use p, when using fdisk. Enter the number of the partition, 1. We now need to enter the amount to allocate to the partition. Do this by entering a + followed by the amount. This is asked in bytes, but its easier to use K, M, G for bigger numbers. Our first partition needs to be 500MB, so we enter +500M. If you did not remove the previous partition you'll get a notification the partion has an ext4 signature. You'll need to enter y to the question if this can be removed.

bash

```
# FDISK
Command (m for help): n
Partition type
   p   primary (0 primary, 0 extended, 4 free)
   e   extended (container for logical partitions)
Select (default p): p
Partition number (1-4, default 1): 1
First sector (2048-16777215, default 2048):
Last sector, +/-sectors or +/-size{K,M,G,T,P} (2048-16777215, default 167

Created a new partition 1 of type 'Linux' and of size 500 MiB.
Partition #1 contains a ext4 signature.

Do you want to remove the signature? [Y]es/[N]o: y

The signature will be removed by a write command.

Command (m for help):
```

bash

```
# GDISK
Command (? for help): n
Partition number (1-128, default 1): 1
First sector (34-16777182, default = 2048) or {+-}size{KMGTP}:
Last sector (2048-16777182, default = 16777182) or {+-}size{KMGTP}: +500M
Current type is 8300 (Linux filesystem)
Hex code or GUID (L to show codes, Enter = 8300):
Changed type of partition to 'Linux filesystem'

Command (? for help):
```

Create the next 2 partitions the same way, sdb2 with 500MB and sdb3 with 300MB.

bash

```
# FDISK
```

```
Command (m for help): n
Partition type
   p   primary (1 primary, 0 extended, 3 free)
   e   extended (container for logical partitions)
Select (default p): p
Partition number (2-4, default 2): 2
First sector (1026048-16777215, default 1026048):
Last sector, +/-sectors or +/-size{K,M,G,T,P} (1026048-16777215, default
```

Created a new partition 2 of type 'Linux' and of size 500 MiB.

```
Command (m for help): n
Partition type
   p   primary (2 primary, 0 extended, 2 free)
   e   extended (container for logical partitions)
Select (default p): p
Partition number (3,4, default 3): 3
First sector (2050048-16777215, default 2050048):
Last sector, +/-sectors or +/-size{K,M,G,T,P} (2050048-16777215, default
```

Created a new partition 3 of type 'Linux' and of size 300 MiB.

```
Command (m for help):
```

bash

```
# GDISK
Command (? for help): n
Partition number (2-128, default 2): 2
First sector (34-16777182, default = 1026048) or {+-}size{KMGTP}:
Last sector (1026048-16777182, default = 16777182) or {+-}size{KMGTP}: +5
Current type is 8300 (Linux filesystem)
Hex code or GUID (L to show codes, Enter = 8300):
Changed type of partition to 'Linux filesystem'
```

```
Command (? for help): n
Partition number (3-128, default 3): 3
First sector (34-16777182, default = 2050048) or {+-}size{KMGTP}:
Last sector (2050048-16777182, default = 16777182) or {+-}size{KMGTP}: +3
```

```
Current type is 8300 (Linux filesystem)
Hex code or GUID (L to show codes, Enter = 8300):
Changed type of partition to 'Linux filesystem'

Command (? for help):
```

For the fourth partition, we change the partition type to extended when using fdisk as mentioned before. With this extended partition we will claim all the free space that is left, so we can create multiple logical partitions within this extended partition (space).

```
bash

# FDISK
Command (m for help): n
Partition type
   p   primary (3 primary, 0 extended, 1 free)
   e   extended (container for logical partitions)
Select (default e): e

Selected partition 4
First sector (2664448-16777215, default 2664448):
Last sector, +/-sectors or +/-size{K,M,G,T,P} (2664448-16777215, default

Created a new partition 4 of type 'Extended' and of size 6.7 GiB.

Command (m for help):
```

We do not make an extended partition with gdisk, we can create up to 128 partitions, so no need to create an extended one! Now create partition 5 with 350MB and 6 with 400MB as before, notice we do not need to choose a partition type this time. This is not possible as all the partition space is used in fdisk. Our fdisk command knows there is an extended partition and that it needs to use this one for any new partitions.

```
bash

# FDISK
Command (m for help): n
```

All primary partitions are **in** use.

Adding logical partition **5**

First sector (**2666496**-16777215, default **2666496**):

Last sector, +/-sectors or +/-size{K,M,G,T,P} (**2666496**-16777215, default

Created a new partition **5** of type '**Linux**' and of size **350** MiB.

Command (m **for** help): n

All primary partitions are **in** use.

Adding logical partition **6**

First sector (**3385344**-16777215, default **3385344**):

Last sector, +/-sectors or +/-size{K,M,G,T,P} (**3385344**-16777215, default

Created a new partition **6** of type '**Linux**' and of size **400** MiB.

Command (m **for** help):

With gdisk we will skip partition 4 to keep the same numbers as with the fdisk demo. Otherwise we wouldn't do that.

bash

GDISK

Command (? **for** help): n

Partition number (4-128, default 4): **5**

First sector (**34**-16777182, default = **2664448**) or {+-}size{KMGTP}:

Last sector (**2664448**-16777182, default = **16777182**) or {+-}size{KMGTP}: +3

Current type is **8300** (Linux filesystem)

Hex code or GUID (L to show codes, Enter = **8300**):

Changed type of partition to '**Linux filesystem**'

Command (? **for** help): n

Partition number (4-128, default 4): **6**

First sector (**34**-16777182, default = **3381248**) or {+-}size{KMGTP}:

Last sector (**3381248**-16777182, default = **16777182**) or {+-}size{KMGTP}: +4

Current type is **8300** (Linux filesystem)

Hex code or GUID (L to show codes, Enter = **8300**):

Changed type of partition to '**Linux filesystem**'

Command (? for help):

Check the partition table by entering p. If we want to check if there is any free space on our drive we can do so by entering F in fdisk, gdisk shows this when pressing p.

bash

FDISK

Command (m for help): p

Disk /dev/sdb: 8 GiB, 8589934592 bytes, 16777216 sectors

Disk model: VMware Virtual S

Units: sectors of 1 * 512 = 512 bytes

Sector size (logical/physical): 512 bytes / 512 bytes

I/O size (minimum/optimal): 512 bytes / 512 bytes

Disklabel type: dos

Disk identifier: 0x8b8ca071

Device	Boot	Start	End	Sectors	Size	Id	Type
/dev/sdb1		2048	1026047	1024000	500M	83	Linux
/dev/sdb2		1026048	2050047	1024000	500M	83	Linux
/dev/sdb3		2050048	2664447	614400	300M	83	Linux
/dev/sdb4		2664448	16777215	14112768	6.7G	5	Extended
/dev/sdb5		2666496	3383295	716800	350M	83	Linux
/dev/sdb6		3385344	4204543	819200	400M	83	Linux

Filesystem/RAID signature on partition 1 will be wiped.

Command (m for help): F

Unpartitioned space /dev/sdb: 5.99 GiB, 6436159488 bytes, 12570624 sector

Units: sectors of 1 * 512 = 512 bytes

Sector size (logical/physical): 512 bytes / 512 bytes

Start	End	Sectors	Size
4206592	16777215	12570624	6G

Command (m for help):

bash

```
# GDISK
Command (? for help): p
Disk /dev/sdb: 16777216 sectors, 8.0 GiB
Model: VMware Virtual S
Sector size (logical/physical): 512/512 bytes
Disk identifier (GUID): 4B226FAA-BFB1-4A6F-9F76-CA70124E6336
Partition table holds up to 128 entries
Main partition table begins at sector 2 and ends at sector 33
First usable sector is 34, last usable sector is 16777182
Partitions will be aligned on 2048-sector boundaries
Total free space is 12578749 sectors (6.0 GiB)
```

Number	Start (sector)	End (sector)	Size	Code	Name
1	2048	1026047	500.0 MiB	8300	Linux filesystem
2	1026048	2050047	500.0 MiB	8300	Linux filesystem
3	2050048	2664447	300.0 MiB	8300	Linux filesystem
5	2664448	3381247	350.0 MiB	8300	Linux filesystem
6	3381248	4200447	400.0 MiB	8300	Linux filesystem

```
Command (? for help):
```

We can see that all partition types are Linux at the moment, we'll change some to swap, FAT32 and Linux LVM. Do this by entering t and the number of the file system we want. To get a list of all file systems enter l, we need to check the number of swap (82), FAT32 (0c) and Linux LVM (8e).

bash

```
# FDISK
Command (m for help): l
```

00 Empty	24 NEC DOS	81 Minix / old Lin	bf Solaris
01 FAT12	27 Hidden NTFS Win	82 Linux swap / So	c1 DRDOS/sec
02 XENIX root	39 Plan 9	83 Linux	c4 DRDOS/sec
03 XENIX usr	3c PartitionMagic	84 OS/2 hidden or	c6 DRDOS/sec
04 FAT16 <32M	40 Venix 80286	85 Linux extended	c7 Syrinx
05 Extended	41 PPC PReP Boot	86 NTFS volume set	da Non-FS dat
06 FAT16	42 SFS	87 NTFS volume set	db CP/M / CT0

07 HPFS/NTFS/exFAT	4d QNX4.x	88 Linux plaintext	de Dell Utili
08 AIX	4e QNX4.x 2nd part	8e Linux LVM	df BootIt
09 AIX bootable	4f QNX4.x 3rd part	93 Amoeba	e1 DOS access
0a OS/2 Boot Manag	50 OnTrack DM	94 Amoeba BBT	e3 DOS R/O
0b W95 FAT32	51 OnTrack DM6 Aux	9f BSD/OS	e4 SpeedStor
0c W95 FAT32 (LBA)	52 CP/M	a0 IBM Thinkpad hi	ea Linux exte
0e W95 FAT16 (LBA)	53 OnTrack DM6 Aux	a5 FreeBSD	eb BeOS fs
0f W95 Ext'd (LBA)	54 OnTrackDM6	a6 OpenBSD	ee GPT
10 OPUS	55 EZ-Drive	a7 NeXTSTEP	ef EFI (FAT-1
11 Hidden FAT12	56 Golden Bow	a8 Darwin UFS	f0 Linux/PA-R
12 Compaq diagnost	5c Priam Edisk	a9 NetBSD	f1 SpeedStor
14 Hidden FAT16 <3	61 SpeedStor	ab Darwin boot	f4 SpeedStor
16 Hidden FAT16	63 GNU HURD or Sys	af HFS / HFS+	f2 DOS second
17 Hidden HPFS/NTF	64 Novell Netware	b7 BSDI fs	fb VMware VMF
18 AST SmartSleep	65 Novell Netware	b8 BSDI swap	fc VMware VMK
1b Hidden W95 FAT3	70 DiskSecure Mult	bb Boot Wizard hid	fd Linux raid
1c Hidden W95 FAT3	75 PC/IX	bc Acronis FAT32 L	fe LANstep
1e Hidden W95 FAT1	80 Old Minix	be Solaris boot	ff BBT

Aliases:

linux	- 83
swap	- 82
extended	- 05
uefi	- EF
raid	- FD
lvm	- 8E
linuxex	- 85

Command (m **for** help): t

Partition number (1-6, default 6): 2

Hex code or alias (type L to list all): 82

Changed type of partition 'Linux' to 'Linux swap / Solaris'.

Command (m **for** help): t

Partition number (1-6, default 6): 5

Hex code or alias (type L to list all): c

Changed type of partition 'Linux' to 'W95 FAT32 (LBA)'.

Command (m **for** help): t

Partition number (1-6, default 6): 6

Hex code or alias (type L to list all): 8e

Changed type of partition 'Linux' to 'Linux LVM'.

bash

GDISK

Command (? **for** help): l

Type search string, or <Enter> to show all codes:

0700 Microsoft basic data	0701 Microsoft Storage Replica
0702 ArcaOS Type 1	0c01 Microsoft reserved
2700 Windows RE	3000 ONIE boot
3001 ONIE config	3900 Plan 9
4100 PowerPC PReP boot	4200 Windows LDM data
4201 Windows LDM metadata	4202 Windows Storage Spaces
7501 IBM GPFS	7f00 ChromeOS kernel
7f01 ChromeOS root	7f02 ChromeOS reserved
8200 Linux swap	8300 Linux filesystem
8301 Linux reserved	8302 Linux /home
8303 Linux x86 root (/)	8304 Linux x86-64 root (/)
8305 Linux ARM64 root (/)	8306 Linux /srv
8307 Linux ARM32 root (/)	8308 Linux dm-crypt
8309 Linux LUKS	830a Linux IA-64 root (/)
830b Linux x86 root verity	830c Linux x86-64 root verity
830d Linux ARM32 root verity	830e Linux ARM64 root verity
830f Linux IA-64 root verity	8310 Linux /var
8311 Linux /var/tmp	8312 Linux user's home
8313 Linux x86 /usr	8314 Linux x86-64 /usr
8315 Linux ARM32 /usr	8316 Linux ARM64 /usr
8317 Linux IA-64 /usr	8318 Linux x86 /usr verity
Press the <Enter> key to see more codes, q to quit:	
8319 Linux x86-64 /usr verity	831a Linux ARM32 /usr verity
831b Linux ARM64 /usr verity	831c Linux IA-64 /usr verity
8400 Intel Rapid Start	8401 SPDK block device
8500 Container Linux /usr	8501 Container Linux resizable r
8502 Container Linux /OEM customization	8503 Container Linux root on RAI
8e00 Linux LVM	a000 Android bootloader

```
a001 Android bootloader 2
a003 Android recovery 1
a005 Android metadata
a007 Android cache
a009 Android persistent
a00b Android fastboot/tertiary
a00d Android vendor
a00f Android factory (alt)
a011 Android EXT
a013 Android SBL2
a015 Android APPSBL
a017 Android QHEE/hyp
a019 Android WDOG debug/sdi
a01b Android CDT
a01d Android SEC
Press the <Enter> key to see more codes, q to quit:
a01f Android misc 1
a021 Android device info
a023 Android MSADP
a025 Android recovery 2
a027 Android modem ST1
a029 Android FSC
a02b Android FSG 2
a02d Android keystore
a02f Android EKSST
a031 Android spare1
a033 Android spare3
a035 Android raw resources
a037 Android FOTA
a039 Android cache
a03b LG (Android) advanced flasher
a03d Android PG2FS
a03f Android MFG
a200 Atari TOS basic data
a501 FreeBSD boot
a503 FreeBSD UFS
a505 FreeBSD Vinum/RAID
Press the <Enter> key to see more codes, q to quit:
a581 Midnight BSD boot
a583 Midnight BSD UFS
a002 Android boot 1
a004 Android misc
a006 Android system 1
a008 Android data
a00a Android factory
a00c Android OEM
a00e Android config
a010 Android meta
a012 Android SBL1
a014 Android SBL3
a016 Android QSEE/tz
a018 Android RPM
a01a Android DDR
a01c Android RAM dump
a01e Android PMIC
a020 Android misc 2
a022 Android APDP
a024 Android DPO
a026 Android persist
a028 Android modem ST2
a02a Android FSG 1
a02c Android SSD
a02e Android encrypt
a030 Android RCT
a032 Android spare2
a034 Android spare4
a036 Android boot 2
a038 Android system 2
a03a Android user data
a03c Android PG1FS
a03e Android board info
a040 Android limits
a500 FreeBSD disklabel
a502 FreeBSD swap
a504 FreeBSD ZFS
a580 Midnight BSD data
a582 Midnight BSD swap
a584 Midnight BSD ZFS
```

a585 Midnight BSD Vinum	a600 OpenBSD disklabel
a800 Apple UFS	a901 NetBSD swap
a902 NetBSD FFS	a903 NetBSD LFS
a904 NetBSD concatenated	a905 NetBSD encrypted
a906 NetBSD RAID	ab00 Recovery HD
af00 Apple HFS/HFS+	af01 Apple RAID
af02 Apple RAID offline	af03 Apple label
af04 AppleTV recovery	af05 Apple Core Storage
af06 Apple SoftRAID Status	af07 Apple SoftRAID Scratch
af08 Apple SoftRAID Volume	af09 Apple SoftRAID Cache
af0a Apple APFS	b300 QNX6 Power-Safe
bb00 Barebox boot loader	bc00 Acronis Secure Zone
be00 Solaris boot	bf00 Solaris root
bf01 Solaris /usr & Mac ZFS	bf02 Solaris swap
bf03 Solaris backup	bf04 Solaris /var
bf05 Solaris /home	bf06 Solaris alternate sector
bf07 Solaris Reserved 1	bf08 Solaris Reserved 2
bf09 Solaris Reserved 3	bf0a Solaris Reserved 4
bf0b Solaris Reserved 5	c001 HP-UX data
Press the <Enter> key to see more codes, q to quit:	
c002 HP-UX service	e100 ONIE boot
e101 ONIE config	e900 Veracrypt data
ea00 XB00TLDR partition	eb00 Haiku BFS
ed00 Sony system partition	ed01 Lenovo system partition
ef00 EFI system partition	ef01 MBR partition scheme
ef02 BIOS boot partition	f800 Ceph OSD
f801 Ceph dm-crypt OSD	f802 Ceph journal
f803 Ceph dm-crypt journal	f804 Ceph disk in creation
f805 Ceph dm-crypt disk in creation	f806 Ceph block
f807 Ceph block DB	f808 Ceph block write-ahead log
f809 Ceph lockbox for dm-crypt keys	f80a Ceph multipath OSD
f80b Ceph multipath journal	f80c Ceph multipath block 1
f80d Ceph multipath block 2	f80e Ceph multipath block DB
f80f Ceph multipath block write-ahead l	f810 Ceph dm-crypt block
f811 Ceph dm-crypt block DB	f812 Ceph dm-crypt block write-a
f813 Ceph dm-crypt LUKS journal	f814 Ceph dm-crypt LUKS block
f815 Ceph dm-crypt LUKS block DB	f816 Ceph dm-crypt LUKS block wr
f817 Ceph dm-crypt LUKS OSD	fb00 VMWare VMFS
fb01 VMWare reserved	fc00 VMWare kcore crash protecti
fd00 Linux RAID	

```

Command (? for help): t
Partition number (1-6): 2
Current type is 8300 (Linux filesystem)
Hex code or GUID (L to show codes, Enter = 8300): L
Type search string, or <Enter> to show all codes: Linux
8200 Linux swap                      8300 Linux filesystem
8301 Linux reserved                  8302 Linux /home
8303 Linux x86 root (/)              8304 Linux x86-64 root (/)
8305 Linux ARM64 root (/)            8306 Linux /srv
8307 Linux ARM32 root (/)            8308 Linux dm-crypt
8309 Linux LUKS                      830a Linux IA-64 root (/)
830b Linux x86 root verity           830c Linux x86-64 root verity
830d Linux ARM32 root verity         830e Linux ARM64 root verity
830f Linux IA-64 root verity         8310 Linux /var
8311 Linux /var/tmp                  8312 Linux user's home
8313 Linux x86 /usr                  8314 Linux x86-64 /usr
8315 Linux ARM32 /usr                8316 Linux ARM64 /usr
8317 Linux IA-64 /usr                8318 Linux x86 /usr verity
8319 Linux x86-64 /usr verity        831a Linux ARM32 /usr verity
831b Linux ARM64 /usr verity         831c Linux IA-64 /usr verity
8500 Container Linux /usr            8501 Container Linux resizable r
8502 Container Linux /OEM customization 8503 Container Linux root on RAI
8e00 Linux LVM                      fd00 Linux RAID
Hex code or GUID (L to show codes, Enter = 8300): 8200
Changed type of partition to 'Linux swap'

```

```

Command (? for help): t
Partition number (1-6): 5
Current type is 8300 (Linux filesystem)
Hex code or GUID (L to show codes, Enter = 8300): EF00
Changed type of partition to 'EFI system partition'

```

```

Command (? for help): t
Partition number (1-6): 6
Current type is 8300 (Linux filesystem)
Hex code or GUID (L to show codes, Enter = 8300): L
Type search string, or <Enter> to show all codes: Linux
8200 Linux swap                      8300 Linux filesystem
8301 Linux reserved                  8302 Linux /home

```

```

8303 Linux x86 root (/)
8305 Linux ARM64 root (/)
8307 Linux ARM32 root (/)
8309 Linux LUKS
830b Linux x86 root verity
830d Linux ARM32 root verity
830f Linux IA-64 root verity
8311 Linux /var/tmp
8313 Linux x86 /usr
8315 Linux ARM32 /usr
8317 Linux IA-64 /usr
8319 Linux x86-64 /usr verity
831b Linux ARM64 /usr verity
8500 Container Linux /usr
8502 Container Linux /OEM customization
8e00 Linux LVM
Hex code or GUID (L to show codes, Enter = 8300): 8e00
Changed type of partition to 'Linux LVM'

Command (? for help):

```

Check all configurations again by entering p. If everything is as expected we can save it by entering w. Again, if this fails, use the partprobe command to save all changes.

```

bash

# FDISK
Command (m for help): p
Disk /dev/sdb: 8 GiB, 8589934592 bytes, 16777216 sectors
Disk model: VMware Virtual S
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disklabel type: dos
Disk identifier: 0x8b8ca071

Device      Boot  Start      End  Sectors  Size Id Type
/dev/sdb1               2048  1026047   1024000  500M 83 Linux

```

```
/dev/sdb2      1026048  2050047  1024000  500M 82 Linux swap / Solaris
/dev/sdb3      2050048  2664447   614400  300M 83 Linux
/dev/sdb4      2664448  16777215 14112768  6.7G  5 Extended
/dev/sdb5      2666496  3383295   716800  350M  c W95 FAT32 (LBA)
/dev/sdb6      3385344  4204543   819200  400M 8e Linux LVM
```

Filesystem/RAID signature on partition **1** will be wiped.

Command (m **for** help): w

The partition table has been altered.

Calling ioctl() to re-read partition table.

Syncing disks.

bash

GDISK

Command (? **for** help): p

Disk /dev/sdb: **16777216** sectors, **8.0** GiB

Model: VMware Virtual S

Sector size (logical/physical): **512/512** bytes

Disk identifier (GUID): 4B226FAA-BFB1-4A6F-9F76-CA70124E6336

Partition table holds up to **128** entries

Main partition table begins at sector **2** and ends at sector **33**

First usable sector is **34**, last usable sector is **16777182**

Partitions will be aligned on **2048**-sector boundaries

Total **free** space is **12578749** sectors (**6.0** GiB)

Number	Start (sector)	End (sector)	Size	Code	Name
1	2048	1026047	500.0 MiB	8300	Linux filesystem
2	1026048	2050047	500.0 MiB	8200	Linux swap
3	2050048	2664447	300.0 MiB	8300	Linux filesystem
5	2664448	3381247	350.0 MiB	EF00	EFI system parti
6	3381248	4200447	400.0 MiB	8E00	Linux LVM

Command (? **for** help): w

Final checks complete. About to **write** GPT data. THIS WILL OVERWRITE EXIST
PARTITIONS!!


```
Do you want to proceed? (Y/N): Y
OK; writing new GUID partition table (GPT) to /dev/sdb.
Warning: The kernel is still using the old partition table.
The new table will be used at the next reboot or after you
run partprobe(8) or kpartx(8)
The operation has completed successfully.
student@linux-ess:~$
```

We are now able to search for all our partitions in the /proc/partitions folder.

```
bash
student@linux-ess:~$ grep sdb /proc/partitions
      8          16    8388608 sdb
      8          17     512000 sdb1
      8          18     512000 sdb2
      8          19    3072000 sdb3
      8          20         1 sdb4
      8          21    3584000 sdb5
      8          22    4096000 sdb6
```

All partitions are now ready with their different types. We now need to enter the command to install the file systems to these partitions. Again we use the mkfs command for sdb1, sdb3 and sdb5. Sdb2 will be a swap space, we use the command mkswap to accomplish this. For the lvm volume, sdb6, we need to create a physical volume with the pvcreate command.

```
bash
student@linux-ess:~$ sudo mkfs.ext4 /dev/sdb1
mke2fs 1.46.5 (30-Dec-2021)
Creating filesystem with 128000 4k blocks and 128000 inodes
Filesystem UUID: 2860cdfd-1eb5-4fad-b5bf-fcd4b9362009
Superblock backups stored on blocks:
    32768, 98304

Allocating group tables: done
Writing inode tables: done
Creating journal (4096 blocks): done
```

Writing superblocks and filesystem accounting information: **done**

```
student@linux-ess:~$ sudo mkswap /dev/sdb2
```

Setting up swapspace version **1**, size = **500** MiB (**524283904** bytes)

no label, **UUID**=c7fc81ae-3382-40ab-b37e-7cdeb76535aa

```
student@linux-ess:~$ sudo mkfs.ext2 /dev/sdb3
```

mke2fs 1.46.5 (**30-Dec-2021**)

Creating filesystem with **76800** 4k blocks and **76800** inodes

Filesystem **UUID**: e9e9e303-c9f1-48a6-97d0-d527b0fb7cd0

Superblock backups stored on blocks:

32768

Allocating group tables: **done**

Writing inode tables: **done**

Writing superblocks and filesystem accounting information: **done**

```
student@linux-ess:~$ sudo mkfs.vfat /dev/sdb5
```

mkfs.fat **4.2** (**2021-01-31**)

```
student@linux-ess:~$ sudo pvcreate /dev/sdb6
```

Physical volume **"/dev/sdb6"** successfully created.

All of the partitions are now ready to be mounted, used as swapspace or added to a volume group. We'll check our partitions and mount point with the fdisk and lsblk commands. We can check all physical volumes with the pvs command.

bash

```
student@linux-ess:~$ sudo fdisk -l
```

...

Disk /dev/sdb: **8** GiB, **8589934592** bytes, **16777216** sectors

Disk model: VMware Virtual S

Units: sectors of **1** * **512** = **512** bytes

Sector size (logical/physical): **512** bytes / **512** bytes

I/O size (minimum/optimal): **512** bytes / **512** bytes

Disklabel type: dos

Disk identifier: 0x8b8ca071

Device	Boot	Start	End	Sectors	Size	Id	Type
/dev/sdb1		2048	1026047	1024000	500M	83	Linux

```

/dev/sdb2      1026048  2050047  1024000  500M 82 Linux swap / Solaris
/dev/sdb3      2050048  2664447   614400  300M 83 Linux
/dev/sdb4      2664448  16777215 14112768   6.7G  5 Extended
/dev/sdb5      2666496  3383295   716800  350M  c W95 FAT32 (LBA)
/dev/sdb6      3385344  4204543   819200  400M 8e Linux LVM
...
student@linux-ess:~$ sudo lsblk -f /dev/sdb
NAME      FSTYPE      FSVER      LABEL UUID
sdb
├─sdb1 ext4          1.0          2860cdfd-1eb5-4fad-b5bf-fcd4b9362009
├─sdb2 swap          1           c7fc81ae-3382-40ab-b37e-7cdeb76535aa
├─sdb3 ext2          1.0          e9e9e303-c9f1-48a6-97d0-d527b0fb7cd0
├─sdb4
├─sdb5 vfat          FAT16        5555-CA20
└─sdb6 LVM2_member LVM2 001      M5f0lN-AczT-hscz-nzs9-iyVe-f49T-V04MNd
student@linux-ess:~$ sudo pvs
PV          VG          Fmt Attr PSize  PFree
/dev/sda3   ubuntu-vg   lvm2 a--   18.22g  8.22g
/dev/sdb6   lvm2       lvm2 ---   400.00m 400.00m

```

We can also check all the block devices with partitions with the parted command.

```

bash

student@linux-ess:~$ sudo parted -l
Model: VMware, VMware Virtual S (scsi)
Disk /dev/sda: 21.5GB
Sector size (logical/physical): 512B/512B
Partition Table: gpt
Disk Flags:

Number  Start   End     Size    File system  Name  Flags
  1      1049kB  2097kB  1049kB             bios_grub
  2      2097kB  1904MB  1902MB   ext4
  3      1904MB  21.5GB  19.6GB

Model: VMware, VMware Virtual S (scsi)

```

Disk /dev/sdb: 8590MB

Sector size (logical/physical): 512B/512B

Partition Table: gpt

Disk Flags:

Number	Start	End	Size	File system	Name	Flag
1	1049kB	525MB	524MB	ext4	Linux filesystem	
2	525MB	1050MB	524MB	linux-swap(v1)	Linux swap	swap
3	1050MB	1364MB	315MB	ext2	Linux filesystem	
5	1364MB	1731MB	367MB		EFI system partition	boot
6	1731MB	2151MB	419MB		Linux LVM	lvm

Model: Linux device-mapper (linear) (dm)

Disk /dev/mapper/myvg0-music: 54.5MB

Sector size (logical/physical): 512B/512B

Partition Table: loop

Disk Flags:

Number	Start	End	Size	File system	Flags
1	0.00B	54.5MB	54.5MB	ext4	

Model: Linux device-mapper (linear) (dm)

Disk /dev/mapper/ubuntu--vg-ubuntu--lv: 10.7GB

Sector size (logical/physical): 512B/512B

Partition Table: loop

Disk Flags:

Number	Start	End	Size	File system	Flags
1	0.00B	10.7GB	10.7GB	ext4	

LVMs

Now that we created a physical volume on our Linux LVM, we'll go a little further into this. Back in the days, when LVM wasn't available, when a drive ran out of free space, we needed to replace it by a larger one and copy everything over

from the old drive. This was time consuming and very inefficient.

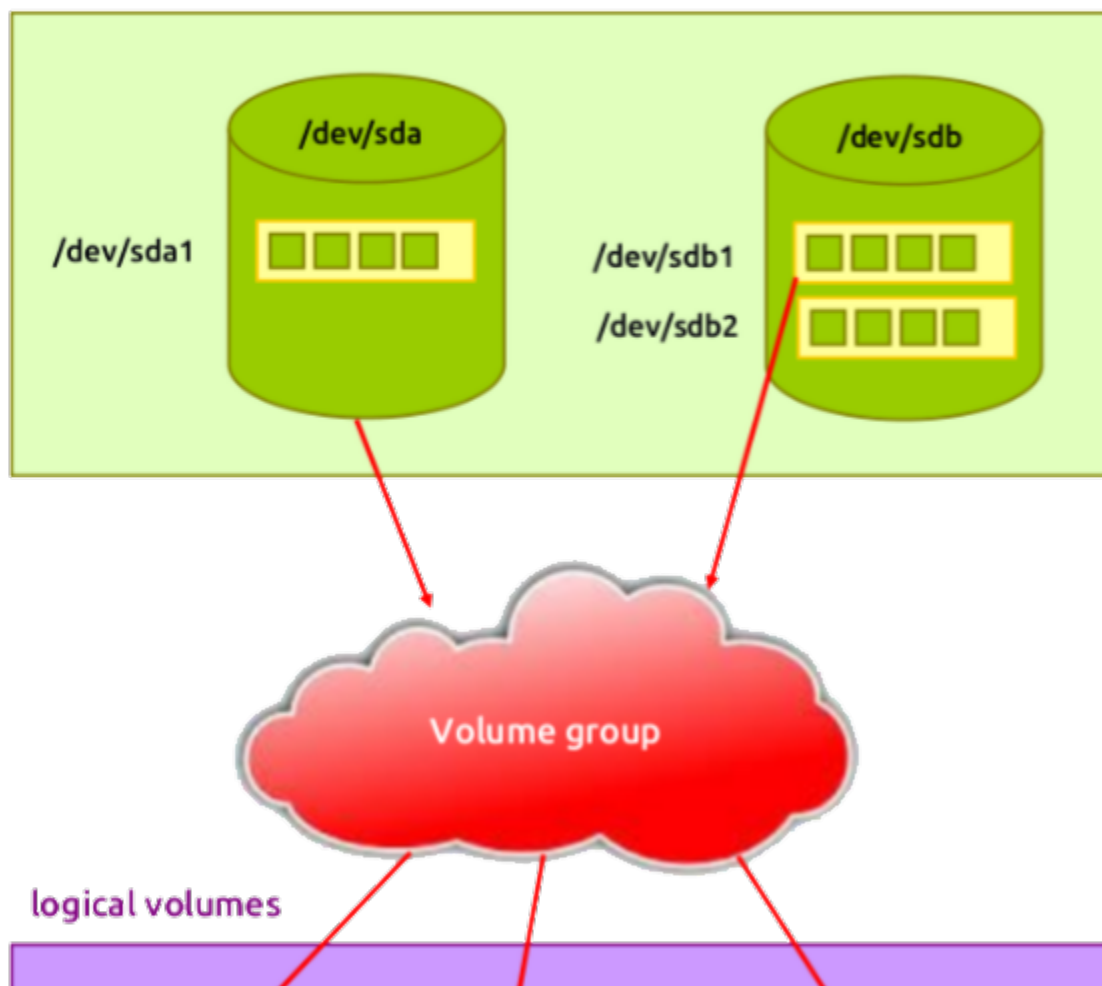
LVM gives us more flexibility to add new physical volumes to a volume group.


This brings a few advantages:

- Adding drive space is possible even when the logical volume is in use
- You can add physical volumes to a volume group whenever needed
- It's possible to move data from one physical volume to another, this way we can replace smaller drives with larger ones without any downtime
- Even downsizing a logical volume is possible if its file system supports it.
- LVM also supports advanced possibilities like mirroring and working with clusters.

LVM OVERVIEW

physical volumes





LVM COMMANDS

Use the command `vgs` to check for volume groups, `pvs` to check for physical volumes and `lvs` to check for logical volumes.

bash

```
student@linux-ess:~$ sudo vgs
VG          #PV #LV #SN Attr   VSize  VFree
ubuntu-vg   1   1   0 wz--n- 18.22g 8.22g
student@linux-ess:~$ sudo pvs
PV          VG          Fmt  Attr PSize   PFree
/dev/sda3   ubuntu-vg  lvm2 a--   18.22g   8.22g
/dev/sdb6                   lvm2 ---   400.00m 400.00m
student@linux-ess:~$ sudo lvs
LV          VG          Attr      LSize   Pool Origin Data%  Meta%  Move  Lc
ubuntu-lv   ubuntu-vg  -wi-ao---- 10.00g
```

First we'll check our main drive, if we have a Linux LVM partition

bash

```
student@linux-ess:~$ sudo fdisk -l /dev/sda
Disk /dev/sda: 20 GiB, 21474836480 bytes, 41943040 sectors
Disk model: VMware Virtual S
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disklabel type: gpt
Disk identifier: A9E92175-1433-43DC-968D-95C5E18A2105
```

Device	Start	End	Sectors	Size	Type
--------	-------	-----	---------	------	------

```

/dev/sda1    2048      4095      2048      1M BIOS boot
/dev/sda2    4096  3719167  3715072  1.8G Linux filesystem
/dev/sda3  3719168 41940991 38221824 18.2G Linux filesystem

```

With the `pvdisk` command we can check if the partition is already in use by a LVM group.

bash

```

student@linux-ess:~$ sudo pvdisk /dev/sda3
--- Physical volume ---
PV Name           /dev/sda3
VG Name           ubuntu-vg
PV Size           <18.23 GiB / not usable 3.00 MiB
Allocatable       yes
PE Size           4.00 MiB
Total PE          4665
Free PE           2105
Allocated PE      2560
PV UUID           9cx3Lm-p95X-Q5f0-ZLS0-KPtf-1ZoN-UyqahD

```

We see that `/dev/sda3` has a physical size of 18.23GiB. This volume is added to the volume group: `ubuntu-vg`. The smallest unit storage that can be granted is 4 MiB.

❗ EXTRA INFO: 1 MB = 1000000 BYTES = 10^6 BYTES
1 MiB = 1048576 BYTES = 2^{20} BYTES

To check the information about this volume group use the command `vgdisplay`.

bash

```

student@linux-ess:~$ sudo vgdisplay
--- Volume group ---
VG Name           ubuntu-vg
System ID
Format            lvm2
Metadata Areas     1

```

```

Metadata Sequence No 2
VG Access              read/write
VG Status              resizable
MAX LV                0
Cur LV               1
Open LV               1
Max PV                0
Cur PV               1
Act PV                1
VG Size               18.22 GiB
PE Size               4.00 MiB
Total PE              4665
Alloc PE / Size       2560 / 10.00 GiB
Free PE / Size        2105 / 8.22 GiB
VG UUID               cPT0cP-GoZh-K91H-rSeQ-2ByY-SCtW-dNqE7C

```

With the `lvdisplay` command, we can check from which volume group it is a member, the total size, etc.

bash

```

student@linux-ess:~$ sudo lvdisplay
--- Logical volume ---
LV Path                /dev/ubuntu-vg/ubuntu-lv
LV Name                ubuntu-lv
VG Name                ubuntu-vg
LV UUID                o7jRC0-02XW-z9dW-t2g7-CQ0o-EtCc-VhxL8v
LV Write Access        read/write
LV Creation host, time ubuntu-server, 2022-07-09 10:36:52 +0000
LV Status               available
# open                 1
LV Size                10.00 GiB
Current LE              2560
Segments               1
Allocation              inherit
Read ahead sectors     auto
- currently set to    256
Block device           253:0

```


We can see that the logical volume `ubuntu-lv` is getting allocation space from the volume group `ubuntu-vg`. In the folder `/dev/mapper` you'll find the file `ubuntu--vg-ubuntu--lv`, this name refers to its logical volume.

bash

```
student@linux-ess:~$ ls /dev/mapper/
control  ubuntu--vg-ubuntu--lv
```

In the `fstab` file we see that, while booting, the following 2 volumes are automatically mounted with their filesystems. The `root`- and `boot`-directory are formatted as `ext4`. We'll now check the connection between all these groups and volumes.

bash

```
student@linux-ess:~$ cat /etc/fstab
# /etc/fstab: static file system information.
#
# Use 'blkid' to print the universally unique identifier for a
# device; this may be used with UUID= as a more robust way to name device
# that works even if disks are added and removed. See fstab(5).
#
# <file system> <mount point> <type> <options> <dump> <pass>
# / was on /dev/ubuntu-vg/ubuntu-lv during curtin installation
/dev/disk/by-id/dm-uuid-LVM-0Iw7F1a0uszZbXwuRqub1PoLZIVWxbIGsWuIcTAED0HcQ
# /boot was on /dev/sda2 during curtin installation
/dev/disk/by-uuid/c23d07d6-5019-4858-8c46-4b31cc88e159 /boot ext4 default
```

```
student@linux-ess:~$ ls -l /dev/ubuntu-vg/*
lrwxrwxrwx 1 root root 7 Sep 15 12:43 /dev/ubuntu-vg/ubuntu-lv -> ../dm-0
student@linux-ess:~$ ls -l /dev/dm-*
brw-rw---- 1 root disk 253, 0 Sep 15 12:43 /dev/dm-0
```

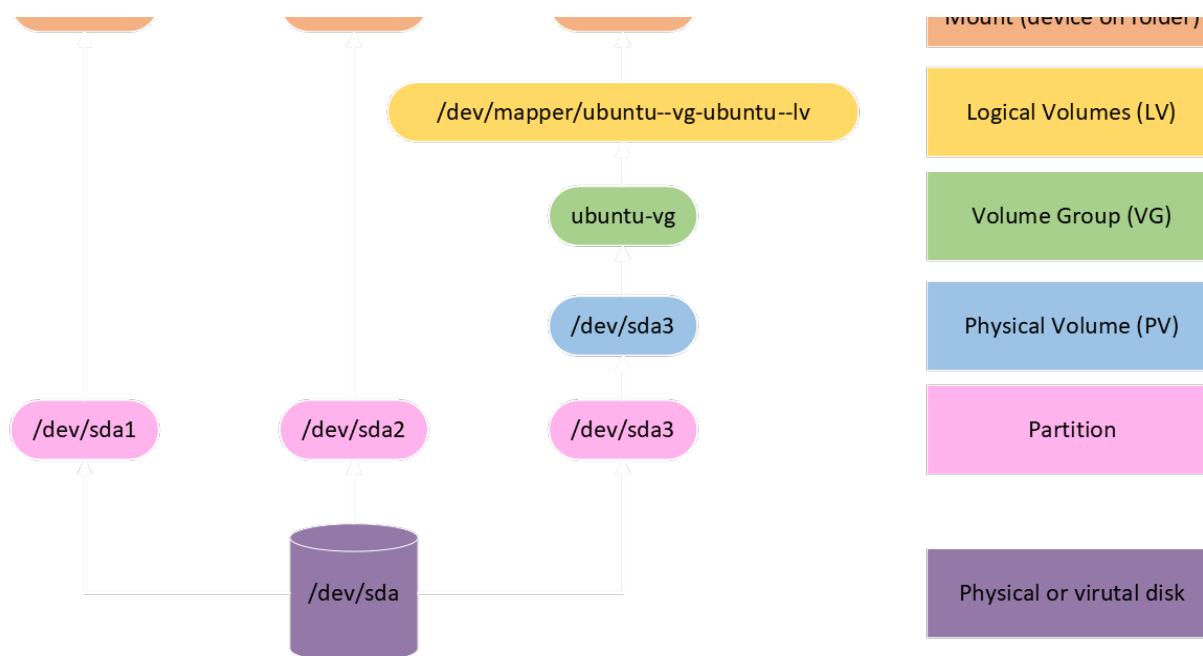
```
student@linux-ess:~$ ls -l /dev/mapper/*
crw----- 1 root root 10, 236 Sep 15 21:08 /dev/mapper/control
lrwxrwxrwx 1 root root 7 Sep 15 21:08 /dev/mapper/ubuntu--vg-ubunt
```

/boot /efi

/boot

/

Filesystem (mkfs device)
Mount (device or folder)



As shown in the figure, it starts with one or more physical volumes, from these you can create a volume group. From this volume group we create logical volumes. We can check these 3 steps with the commands used earlier: `pvdisk`, `vgdisplay` and `lvdisplay`. We will now create our own volume group and logical volumes from the physical volume created earlier. To create this physical volume we used the `pvcreate` command, to create a volume group we can use a similar command: `vgcreate`, we name our group `myvg0`.

bash

```
student@linux-ess:~$ sudo vgcreate myvg0 /dev/sdb6
[sudo] password for student:
Volume group "myvg0" successfully created
student@linux-ess:~$ sudo vgdisplay myvg0
sudo vgdisplay myvg0
--- Volume group ---
VG Name                myvg0
System ID
Format                  lvm2
Metadata Areas          1
Metadata Sequence No    1
VG Access                read/write
VG Status                resizable
MAX LV                  0
Cur LV                  0
```

```

Open LV          0
Max PV           0
Cur PV          1
Act PV           1
VG Size          396.00 MiB
PE Size          4.00 MiB
Total PE         99
Alloc PE / Size  0 / 0
Free PE / Size   99 / 396.00 MiB
VG UUID          AtpRkW-tbTs-MSIv-mXlY-mgv0-H1bJ-gD7ryc

```

From the free 400MiB, 396 MiB is usable in blocks of 4MiB. We'll now create a logical volume 'music' with our group with the lvcreate command.

bash

```

student@linux-ess:~$ sudo lvcreate -n music -L 100M myvg0
Logical volume "music" created.
student@linux-ess:~$ sudo lvs myvg0
LV      VG      Attr      LSize   Pool Origin Data%  Meta%  Move Log Cpy%S
music  myvg0  -wi-a----- 100.00m
student@linux-ess:~$ ls /dev/mapper/myvg0*
/dev/mapper/myvg0-music

```

If we want to use this logical volume, we need to give it a file system and mount it to our system. This is possible, as shown before, with the mount command or an entry in the /etc/fstab file.

bash

```

student@linux-ess:~$ sudo mkfs.ext4 /dev/mapper/myvg0-music
mke2fs 1.46.5 (30-Dec-2021)
Creating filesystem with 25600 4k blocks and 25600 inodes

Allocating group tables: done
Writing inode tables: done
Creating journal (1024 blocks): done
Writing superblocks and filesystem accounting information: done

```

```

student@linux-ess:~$ sudo mkdir /mnt/mymusic
student@linux-ess:~$ sudo mount /dev/mapper/myvg0-music /mnt/mymusic/
student@linux-ess:~$ sudo df -h /mnt/mymusic/
Filesystem                Size  Used Avail Use% Mounted on
/dev/mapper/myvg0-music   90M   24K   83M   1% /mnt/mymusic
student@linux-ess:~$ sudo nano /etc/fstab
/dev/mapper/myvg0-music /mnt/mymusic ext4 default 1 2

```

If we now want to enlarge our logical volume, we can do so without unmounting it. We do need to have free volume in our volume group. First you'll need to enlarge the logical volume and afterwards the file system. So first, we'll check the available space in our group and logical volume.

bash

```

student@linux-ess:~$ sudo df -h /mnt/mymusic/
Filesystem                Size  Used Avail Use% Mounted on
/dev/mapper/myvg0-music   90M   24K   83M   1% /mnt/mymusic
student@linux-ess:~$ sudo vgs myvg0
VG      #PV #LV #SN Attr   VSize   VFree
myvg0   1   1   0 wz--n- 396.00m 296.00m

```

We can then extend the volume as possible with the `lvextend` command.

bash

```

student@linux-ess:~$ sudo lvextend -L +100M /dev/mapper/myvg0-music
Size of logical volume myvg0/music changed from 100.00 MiB (25 extents)
Logical volume myvg0/music successfully resized.

```

Afterwards the file system needs to be enlarged as well, this is possible with the `resize2fs` command for ext filesystems and `xfs_growfs` for xfs filesystems. To do it in just one step you could have given the `-r` option within the previous `lvextend` command.

bash

```

student@linux-ess:~$ sudo resize2fs /dev/mapper/myvg0-music
resize2fs 1.46.5 (30-Dec-2021)

```

```
Filesystem at /dev/mapper/myvg0-music is mounted on /mnt/mymusic; on-line
old_desc_blocks = 1, new_desc_blocks = 1
The filesystem on /dev/mapper/myvg0-music is now 51200 (4k) blocks long.
```

Check if the file system is enlarged after the configuration.

bash

```
student@linux-ess:~$ sudo df -h /mnt/mymusic/
Filesystem                Size  Used Avail Use% Mounted on
/dev/mapper/myvg0-music  184M  120K  175M   1% /mnt/mymusic
```

Now that we have seen how to enlarge a volume, we can try and make it smaller again as well. To start this we need to unmount our volume and check the logical volume for defragmentation with the e2fsck command. Next we need to resize the file system first to the newly wanted smaller size. Afterwards we can make the volume group smaller. Now we can remount the logical volume and check if all configurations were successful.

bash

```
student@linux-ess:~$ student@linux-ess:~$ sudo umount /mnt/mymusic
student@linux-ess:~$ sudo e2fsck -f /dev/mapper/myvg0-music
e2fsck 1.46.5 (30-Dec-2021)
Pass 1: Checking inodes, blocks, and sizes
Pass 2: Checking directory structure
Pass 3: Checking directory connectivity
Pass 4: Checking reference counts
Pass 5: Checking group summary information
/dev/mapper/myvg0-music: 11/51200 files (9.1% non-contiguous), 4262/51200
student@linux-ess:~$ sudo resize2fs /dev/mapper/myvg0-music 50M
resize2fs 1.46.5 (30-Dec-2021)
Resizing the filesystem on /dev/mapper/myvg0-music to 12800 (4k) blocks.
The filesystem on /dev/mapper/myvg0-music is now 12800 (4k) blocks long.

student@linux-ess:~$ sudo lvreduce -L 50M -r /dev/mapper/myvg0-music
Rounding size to boundary between physical extents: 52.00 MiB.
fsck from util-linux 2.37.2
/dev/mapper/myvg0-music: clean, 11/25600 files, 2646/12800 blocks
```

```
resize2fs 1.46.5 (30-Dec-2021)
```

```
Resizing the filesystem on /dev/mapper/myvg0-music to 13312 (4k) blocks.
```

```
The filesystem on /dev/mapper/myvg0-music is now 13312 (4k) blocks long.
```

```
Size of logical volume myvg0/music changed from 200.00 MiB (50 extents)
```

```
Logical volume myvg0/music successfully resized.
```

```
student@linux-ess:~$ sudo mount /dev/mapper/myvg0-music /mnt/mymusic/
```

```
student@linux-ess:~$ sudo df -h /mnt/mymusic/
```

Filesystem	Size	Used	Avail	Use%	Mounted on
/dev/mapper/myvg0-music	42M	72K	39M	1%	/mnt/mymusic

The mount command

We've now mostly mounted our volumes and drives with the mount command and shown off the /etc/fstab file. But the mount command isn't only used to mount local storage devices. We can also mount network directories with NFS or Samba, mount image files and mount drives or USB flash drives which didn't automount. To check what file systems are supported by the kernel to mount, check the /proc/filesystems file. These file systems are supported by the kernel, but not necessarily by your Linux distribution!

bash

```
student@linux-ess:~$ cat /proc/filesystems
```

```
nodev  sysfs
```

```
nodev  tmpfs
```

```
nodev  bdev
```

```
nodev  proc
```

```
nodev  cgroup
```

```
nodev  cgroup2
```

```
nodev  cpuset
```

```
nodev  devtmpfs
```

```
nodev  configfs
```

```
nodev  debugfs
```

```
nodev  tracefs
```

```
nodev  securityfs
```

```
nodev  sockfs
```

```
nodev  bpf
```

```

nodev pipefs
nodev ramfs
nodev hugetlbfs
nodev devpts
    ext3
    ext2
    ext4
    squashfs
    vfat
nodev ecryptfs
    fuseblk
nodev fuse
nodev fusectl
nodev mqueue
nodev pstore
    btrfs
nodev autofs

```

To check all modules which can be loaded by your kernel when a file system is mounted use the command:

```

bash
student@linux-ess:~$ ls /lib/modules/$(uname -r)/kernel/fs
9p      befs      ceph      efs      freevxfs  hfsplus  ksmbd     nfs_comr
adfs     bfs       cifs      erofs    fscache   hpfs     lockd     nfsd
affs     binfmt_misc.ko  coda     exfat    fuse      isofs    minix     nilfs2
afs      btrfs     cramfs    f2fs     gfs2      jffs2    netfs     nls
autofs   cachefiles  dlm      fat      hfs       jfs      nfs       ntfs

```

For more information check the manpage on filesystems: `man fs`

Swapspace

Up next, we check how to use a swapspace. We already created one before, but did not add it to our usable swapspace. The swapspace is used when your system runs out of RAM and needs to offload any unused data, which is needed

later on again. To create a swapspace we use the command `mkswap`, to turn the swapspace on or off, use the `swapon` or `swapoff` command. We'll recreate adding swapspace by the following example. First we check the available space at this time.

```
bash
student@linux-ess:~$ free -h
```

	total	used	free	shared	buff/cache	av
Mem:	3.8Gi	304Mi	2.8Gi	1.0Mi	678Mi	
Swap:	0B	0B	0B			

We'll now create a file of 1 GiB and make a swapspace out of this. At the end we turn it on to make it usable for our system. We can check this by just entering the `swapon` command without parameters. We can also check the full available swapspace as shown earlier.

```
bash
student@linux-ess:~$ sudo dd if=/dev/zero of=/var/opt/myswap bs=1M count=
1024+0 records in
1024+0 records out
1073741824 bytes (1.1 GB, 1.0 GiB) copied, 4.21084 s, 255 MB/s
student@linux-ess:~$ sudo chmod 0600 /var/opt/myswap
student@linux-ess:~$ sudo mkswap /var/opt/myswap
Setting up swapspace version 1, size = 1024 MiB (1073737728 bytes)
no label, UUID=0a79aa31-fca6-4ad7-893b-4d3da75233fb
student@linux-ess:~$ sudo swapon /var/opt/myswap
student@linux-ess:~$ swapon
```

NAME	TYPE	SIZE	USED	PRI0
/var/opt/myswap	file	1024M	0B	-2

```
student@linux-ess:~$ free -h
```

	total	used	free	shared	buff/cache	av
Mem:	3.8Gi	308Mi	1.8Gi	1.0Mi	1.7Gi	
Swap:	1.0Gi	0B	1.0Gi			

Again we can add this to the `/etc/fstab` file to make this extra swapspace permanently available.

bash

```
student@linux-ess:~$ sudo nano /etc/fstab
/var/opt/myswap swap swap defaults 0 0
```

We use swap in the second field as there is no mounting point for swapspace. To turn on any swapspace that is mentioned in the /etc/fstab file, for example with the option noauto, immediately we can use the swapon -a command. If we want to remove swapspace we first need to check that it isn't in use. If this is the case we can unmount it by the swapoff command.

bash

```
student@linux-ess:~$ swapon
NAME                TYPE  SIZE USED PRIO
/var/opt/myswap file 1024M  0B   -2
student@linux-ess:~$ sudo swapoff /var/opt/myswap
```

[< Previous](#)

11 Managing running processes

[Next >](#)

Lab